

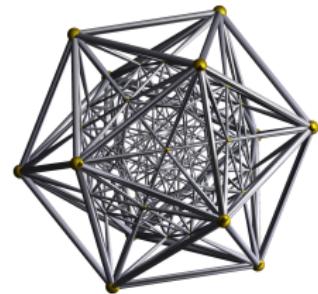
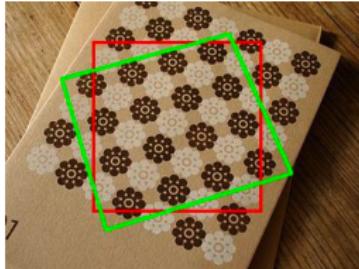
Computational Principles for High-dim Data Analysis

(Lecture Twenty)

Yi Ma

EECS Department, UC Berkeley

November 9, 2021



Structured Nonlinear Low-Dimensional Models

Transform Invariant/Equivariant Low-Rank Texture

- ① Low-Rank Models for Images
- ② Image Inpainting as Matrix Completion
- ③ Transform Invariant Low-Rank Textures
- ④ Applications of TILT

“What humans do with the language of mathematics is to describe patterns.”

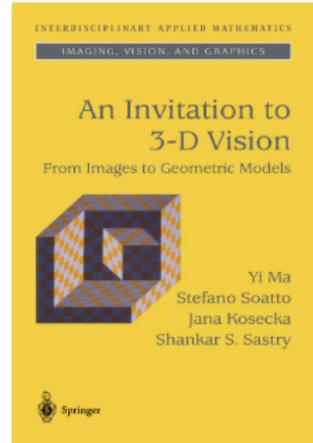
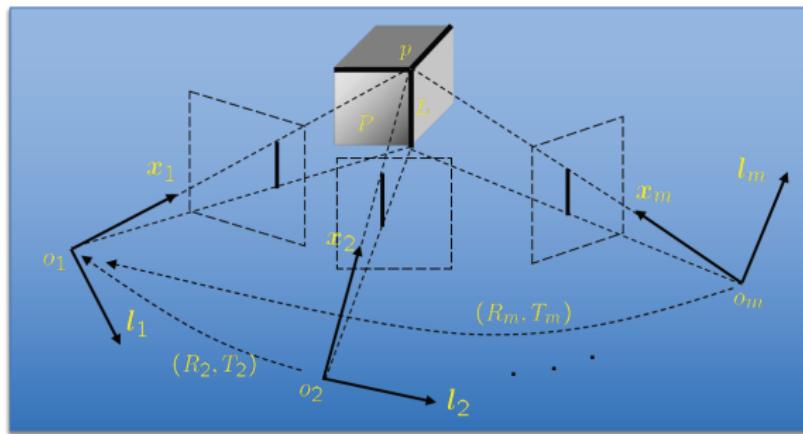
– Lynn A. Steen

Importance of Mathematical Modeling

**If you formulate a problem correctly,
you have more than halfway solved it!**

Example: Rank Conditions for Multiple-View Geometry

Geometric and algorithmic foundations for 3D Reconstruction from images.

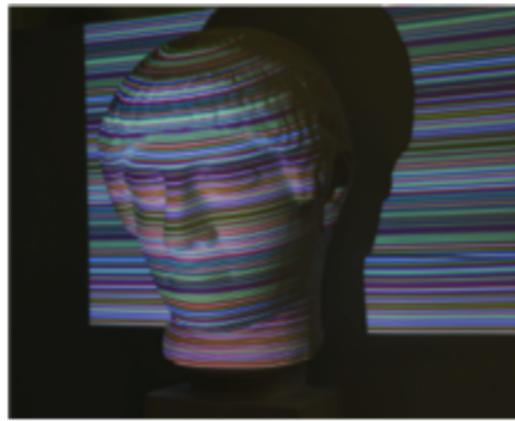


All multi-view incidence conditions among points, lines, planes and symmetric objects are captured by the same rank condition (Ma, 2003):

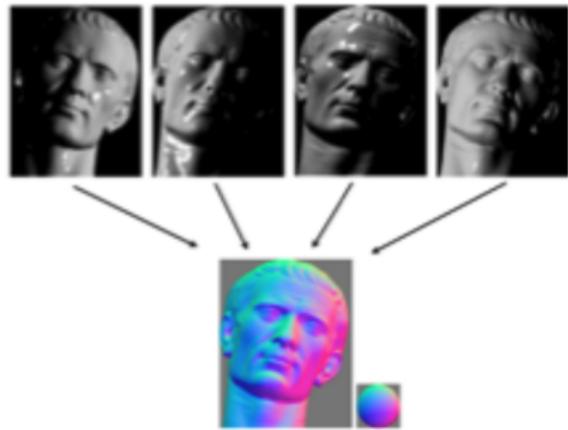
$$\text{rank}(M) = 1 \text{ (or 2)}. \quad (1)$$

Example: Rank Conditions for Photometric Stereo

Geometric and algorithmic foundations for 3D Reconstruction from images.



(b) structured lights



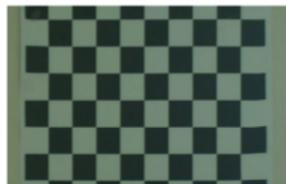
(c) photometric stereo

Images of a (Lambertian) object under different lighting conditions, viewed as columns of a matrix, always satisfy a rank condition ([Chapter 14](#)):

$$\text{rank}(M) = 3. \quad (2)$$

Example: Low-Rank Structures in a Single Image

Low-dimensional structures arise at all spatial scales even in an individual image, especially that of man-made objects.



(a) a calibration rig



(b) a carpet



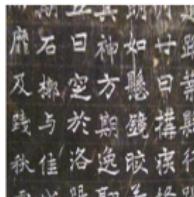
(c) windows



(d) a door



(e) a license plate



(f) characters



(g) a car



(h) a face

Be aware though: only at a rectified view and without any occlusion!

Mathematical Model for Low-Rank Image (Regions)

An image (region) $\mathbf{I}_o(x, y)$ can be explicitly factorized as the combination of r rank-1 functions:

$$\mathbf{I}_o(x, y) \doteq \sum_{i=1}^r u_i(x) \cdot v_i(y). \quad (3)$$

Symmetry: two low-rank textures *equivalent* if they are scaled and translated versions of each other:

$$\mathbf{I}_o(x, y) \sim \alpha \cdot \mathbf{I}_o(ax + t_1, by + t_2),$$

for some $\alpha, a, b, c \in \mathbb{R}_+, t_1, t_2 \in \mathbb{R}$.

Image Inpainting as Matrix Completion

Example: Low-rank image inpainting as matrix completion:

$$\min_{\mathbf{L}} \text{rank}(\mathbf{L}) \quad \text{s.t.} \quad \mathbf{L}(i, j) = \mathbf{I}_o(i, j) \quad \forall (i, j) \in \Omega. \quad (4)$$

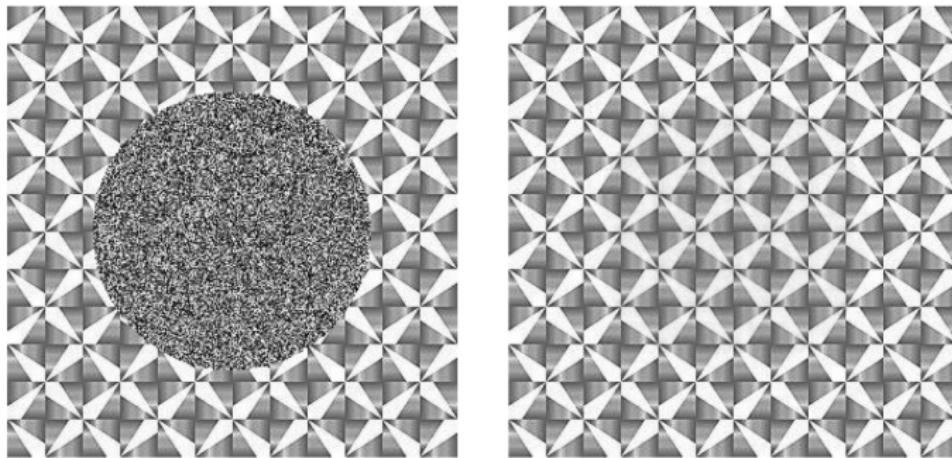
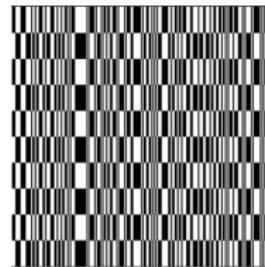
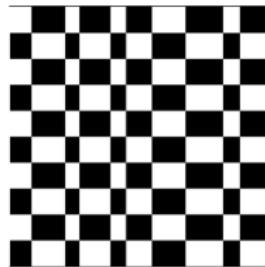
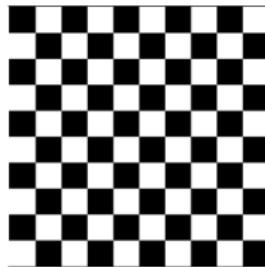


Image Inpainting as Matrix Completion



Is low-rank the **only** low-dim structure in such images?

$$\mathbf{I}_o(x, y) = \sum_{i=1}^r u_i(x)v_i(y) = \mathbf{U}\mathbf{V}^*.$$

\mathbf{U} and \mathbf{V} might themselves be **sparse** in some bases

$\mathbf{U} = \mathbf{B}_1\mathbf{X}_1, \mathbf{V} = \mathbf{B}_2\mathbf{X}_2:$

$$\mathbf{I}_o = \mathbf{B}_1\mathbf{X}_1\mathbf{X}_2^*\mathbf{B}_2^* \doteq \mathbf{B}_1\mathbf{W}_o\mathbf{B}_2^*, \text{ with } \mathbf{W}_o = \mathbf{X}_1\mathbf{X}_2^* \text{ sparse.}$$

Image Inpainting as Matrix Completion

Impose **low-rank and sparse** structures simultaneously (not RPCA!):

$$\min_{\mathbf{L}, \mathbf{W}} \text{rank}(\mathbf{L}) + \lambda \|\mathbf{W}\|_0 \quad \text{s.t.} \quad \mathcal{P}_\Omega[\mathbf{L}] = \mathcal{P}_\Omega[\mathbf{I}], \quad \mathbf{L} = \mathbf{B}_1 \mathbf{W} \mathbf{B}_2^*. \quad (5)$$

Relaxation to a corresponding convex program (why?):

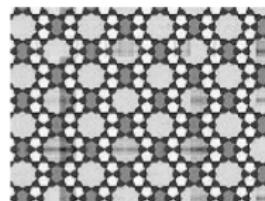
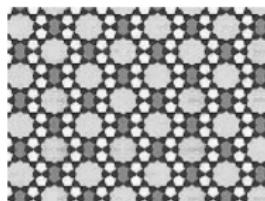
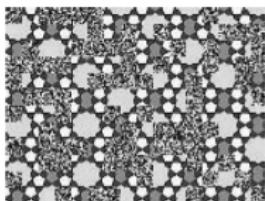
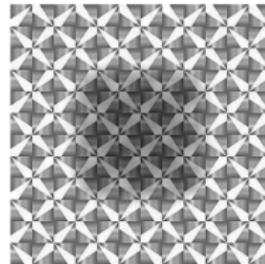
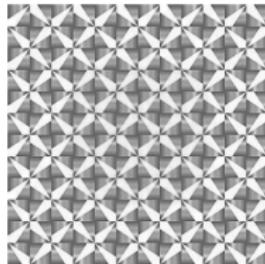
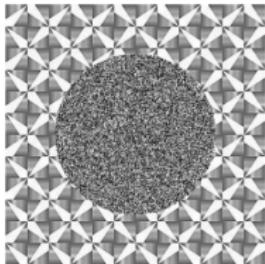
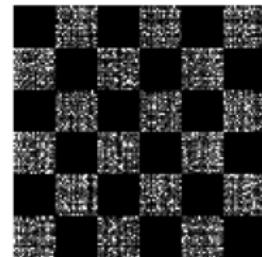
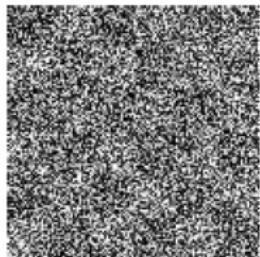
$$\min_{\mathbf{W}} \|\mathbf{W}\|_* + \lambda \|\mathbf{W}\|_1 \quad \text{s.t.} \quad \mathcal{P}_\Omega[\mathbf{B}_1 \mathbf{W} \mathbf{B}_2^*] = \mathcal{P}_\Omega[\mathbf{I}]. \quad (6)$$

Reformulation for a better optimization algorithm:

$$\min_{\mathbf{L}, \mathbf{W}} \|\mathbf{L}\|_* + \lambda \|\mathbf{W}\|_1 \quad \text{s.t.} \quad \mathbf{L} = \mathbf{W}, \quad \mathcal{P}_\Omega[\mathbf{B}_1 \mathbf{W} \mathbf{B}_2^*] = \mathcal{P}_\Omega[\mathbf{I}]. \quad (7)$$

Note that the above problem can be solved by ALM + ADMM!

Image Inpainting: Improved Performance



(a) Input image

(b) Sparse + Low-rank

(c) Low-rank only

Image Recovery

Robustness to corruption $\mathbf{I} = \mathbf{I}_o + \mathbf{E}_o$ for unknown sparse \mathbf{E}_o :

$$\min_{\mathbf{W}} \|\mathbf{W}\|_* + \lambda \|\mathbf{W}\|_1 + \alpha \|\mathbf{E}\|_1 \quad \text{s.t.} \quad \mathcal{P}_{\Omega}[\mathbf{B}_1 \mathbf{W} \mathbf{B}_2^* + \mathbf{E}] = \mathcal{P}_{\Omega}[\mathbf{I}]. \quad (8)$$

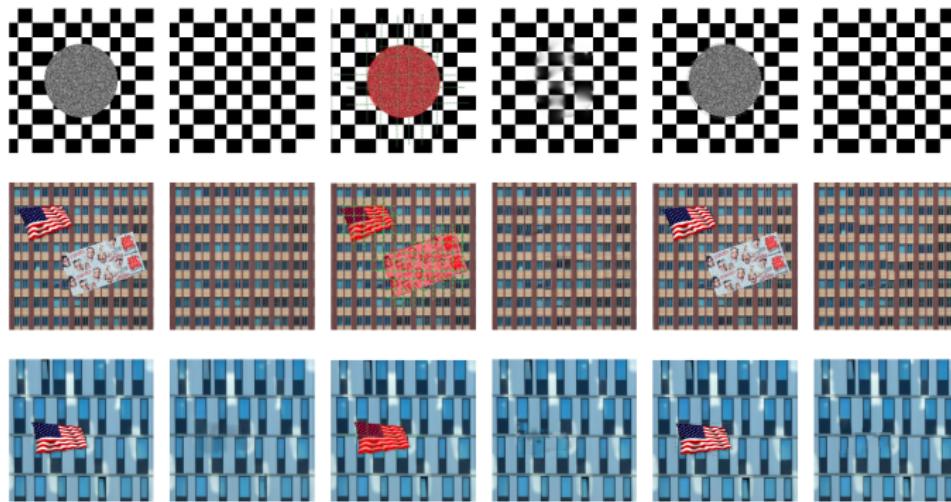


Figure: Left: Low-rank; Middle: Microsoft; Right: Adobe.

Deformation and Corruption

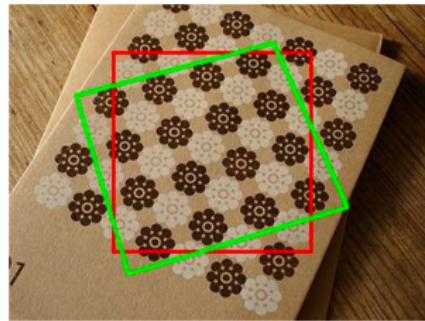
In practice, the observed image \mathbf{I} is typically a **transformed** version of the (rectified) low-rank image \mathbf{I}_o : $\mathbf{I} \circ \tau(x, y) = \mathbf{I}_o(x, y)$ or:

$$\mathbf{I}(x, y) = \mathbf{I}_o \circ \tau^{-1}(x, y) = \mathbf{I}_o(\tau^{-1}(x, y)), \quad \tau \in \mathbb{G}.$$



(a) Low-rank texture \mathbf{I}_o

$$\xleftarrow{\tau}$$



(b) Its image \mathbf{I} under a different viewpoint

Note: the observed image might also be corrupted: $\mathbf{I} = (\mathbf{I}_o + \mathbf{E}_o) \circ \tau^{-1}$.

Algorithm Development

Original problem: transformed low-rank and sparse decomposition

$$\min_{\mathbf{L}, \mathbf{E}, \tau} \text{rank}(\mathbf{L}) + \gamma \|\mathbf{E}\|_0 \quad \text{subject to} \quad \mathbf{I} \circ \tau = \mathbf{L} + \mathbf{E}. \quad (9)$$

Relaxation: use convex surrogates for rank and sparsity

$$\min_{\mathbf{L}, \mathbf{E}, \tau} \underbrace{\|\mathbf{L}\|_* + \lambda \|\mathbf{E}\|_1}_{\text{convex}} \quad \text{subject to} \quad \underbrace{\mathbf{I} \circ \tau = \mathbf{L} + \mathbf{E}}_{\text{still nonlinear!}}. \quad (10)$$

Linearization to deal with nonlinearity: $\mathbf{I} \circ \tau + \nabla \mathbf{I} \cdot d\tau \approx \mathbf{L} + \mathbf{E}$ and

$$\min_{\mathbf{L}, \mathbf{E}, d\tau} \|\mathbf{L}\|_* + \lambda \|\mathbf{E}\|_1 \quad \text{subject to} \quad \mathbf{I} \circ \tau + \nabla \mathbf{I} \cdot d\tau = \mathbf{L} + \mathbf{E}. \quad (11)$$

Note: this is exactly a **Compressive PCP** problem (in Chapter 5)!

The TILT Algorithm

INPUT: Input image $\mathbf{I} \in \mathbb{R}^{w \times h}$, initial transformation $\tau \in \mathbb{G}$ (affine or projective), and a weight $\lambda > 0$.

WHILE not converged **DO**

Step 1: Normalization and compute Jacobian:

$$\mathbf{I} \circ \tau \leftarrow \frac{\mathbf{I} \circ \tau}{\|\mathbf{I} \circ \tau\|_F}; \quad \nabla \mathbf{I} \leftarrow \frac{\partial}{\partial \zeta} \left(\frac{\text{vec}(\mathbf{I} \circ \zeta)}{\|\text{vec}(\mathbf{I} \circ \zeta)\|_2} \right) \Big|_{\zeta=\tau}; \quad (12)$$

Step 2 (inner loop): Solve the linearized **CPCP** problem:

$$\begin{aligned} (\mathbf{L}_*, \mathbf{E}_*, d\tau_*) &\leftarrow \arg \min_{\mathbf{L}, \mathbf{E}, d\tau} \|\mathbf{L}\|_* + \lambda \|\mathbf{E}\|_1 \\ \text{subject to } \mathbf{I} \circ \tau + \nabla \mathbf{I} \cdot d\tau &= \mathbf{L} + \mathbf{E}; \end{aligned} \quad (13)$$

Step 3: Update the transformation: $\tau \leftarrow \tau + d\tau_*$;

END WHILE

OUTPUT: Converged solution \mathbf{L}_* , \mathbf{E}_* , τ_* to problem (10).

Inner Loop of TILT

Apply ALM + ADMM to solve the inner loop CPCP problem (13):

INPUT: The current (deformed and normalized) image $\mathbf{I} \circ \tau \in \mathbb{R}^{m \times n}$ and its Jacobian $\nabla \mathbf{I}$ against current deformation τ (from the outer loop), and $\lambda > 0$.

Initialization: $k = 0, \mathbf{Y}_0 = 0, \mathbf{E}_0 = 0, d\tau_0 = 0, \mu_0 > 0, \rho > 1$;

WHILE not converged **DO**

$$(\mathbf{U}_k, \boldsymbol{\Sigma}_k, \mathbf{V}_k) = \text{SVD}(\mathbf{I} \circ \tau + \nabla \mathbf{I} \cdot d\tau_k - \mathbf{E}_k + \mu_k^{-1} \mathbf{Y}_k);$$

$$\mathbf{L}_{k+1} = \mathbf{U}_k \text{soft}(\boldsymbol{\Sigma}_k, \mu_k^{-1}) \mathbf{V}_k^*;$$

$$\mathbf{E}_{k+1} = \text{soft}(\mathbf{I} \circ \tau + \nabla \mathbf{I} \cdot d\tau_k - \mathbf{L}_{k+1} + \mu_k^{-1} \mathbf{Y}_k, \lambda \mu_k^{-1});$$

$$d\tau_{k+1} = (\nabla \mathbf{I})^\dagger(-\mathbf{I} \circ \tau + \mathbf{L}_{k+1} + \mathbf{E}_{k+1} - \mu_k^{-1} \mathbf{Y}_k);$$

$$\mathbf{Y}_{k+1} = \mathbf{Y}_k + \mu_k(\mathbf{I} \circ \tau + \nabla \mathbf{I} \cdot d\tau_{k+1} - \mathbf{L}_{k+1} - \mathbf{E}_{k+1});$$

$$\mu_{k+1} = \rho \mu_k;$$

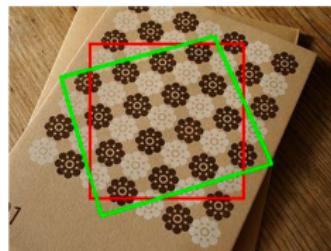
END WHILE

OUTPUT: Converged solution $(\mathbf{L}_*, \mathbf{E}_*, d\tau_*)$ to problem (11).

Rectifying Planar Low-Rank Textures

(a) Low-rank texture \mathbf{I}_o

$$\xleftarrow{\tau}$$

(b) Its image \mathbf{I} under a different viewpoint

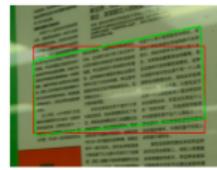
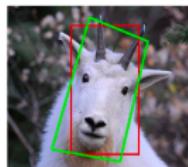
\mathbf{I}_o is on a planar surface and \mathbf{I} is seen from a different view:

$$\mathbf{I}(x, y) = \mathbf{I}_o \circ \tau^{-1}(x, y) = \mathbf{I}_o(\tau^{-1}(x, y))$$

where τ is any planar homography $\mathbf{H} \in \text{GL}(3)$:

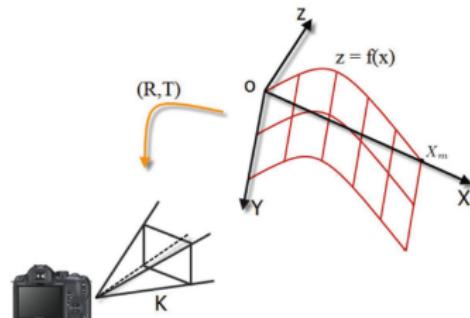
$$\tau(x, y) = \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \sim \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}. \quad (14)$$

Empirical Observations: large range of attraction



Rectifying Generalized Cylinder Surfaces

In practice, very often the low-rank texture I_o does not lie on a planar surface, instead say on a cylindrical surface:



The overall deformation consists of composition of **three** mappings:

- from a flat 2D plane to the 3D cylindrical surface (shape);
- from the surface to the image plane (camera pose and projection);
- from the image plane to the pixel coordinates (intrinsic calibration).

Rectifying Generalized Cylinder Surfaces

Then $\mathbf{I}_o = \mathbf{I} \circ g$ where the transformation g from the texture coordinates $\mathbf{I}_o(x, y)$ to the image coordinates $\mathbf{I}(u, v)$ is a composition of the three mappings specified above:

$$g : (x, y) \mapsto (X_c, Y_c, Z_c) \mapsto (x_n, y_n) \mapsto (u, v). \quad (15)$$

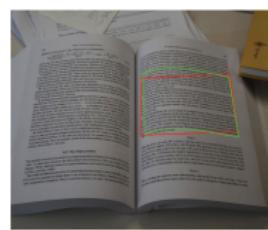
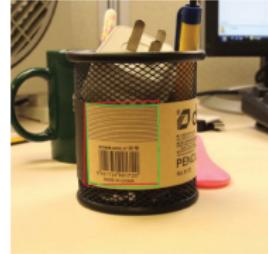
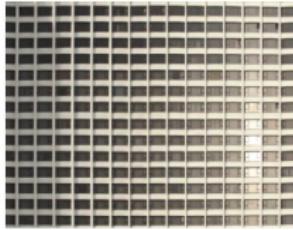
Hence we solve the following **nonlinear** optimization problem:

$$\min_{\mathbf{L}, \mathbf{E}, c, \mathbf{R}, \mathbf{T}} \|\mathbf{L}\|_* + \lambda \|\mathbf{E}\|_1 \quad \text{s.t.} \quad \mathbf{I} \circ g = \mathbf{L} + \mathbf{E}. \quad (16)$$

As in TILT, this can be solved **iteratively** from its linearized version:

$$\min_{\mathbf{L}, \mathbf{E}, dg} \|\mathbf{L}\|_* + \lambda \|\mathbf{E}\|_1 \quad \text{s.t.} \quad \mathbf{I} \circ g + \nabla \mathbf{I}_g \cdot dg = \mathbf{L} + \mathbf{E}, \quad (17)$$

Empirical Results



Camera Calibration

Three most important factors in robotics and augmented reality:
calibration, calibration, calibration!



Figure 15.10 Left: typical image of a fisheye camera. Right: image of a perspective camera.



Figure 15.11 Left two: Images of a typical calibration rig. Right: Corners need to be marked (or detected) for conventional calibration methods.

Camera Calibration: from a calibration rig to the image

Extrinsic parameters τ_i (camera pose – different for each image i):

rotation $\mathbf{R} \in SO(3)$ and translation $\mathbf{T} \in \mathbb{R}^3$.

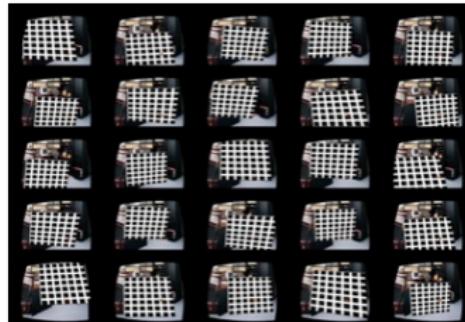
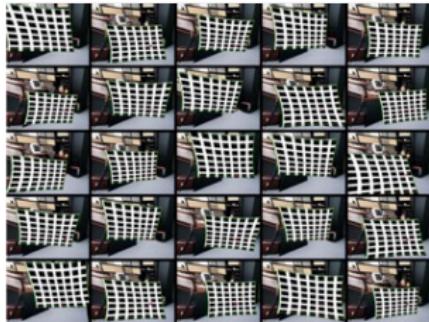
Intrinsic parameters τ_0 (focal length etc. – common for all images):

$$\mathbf{K} \doteq \begin{bmatrix} f_x & \theta & o_x \\ 0 & f_y & o_y \\ 0 & 0 & 1 \end{bmatrix} \in \mathbb{R}^{3 \times 3}. \quad (18)$$

Lens distortion τ_0 (radial distortion etc. – common for all images):

$$\left\{ \begin{array}{lcl} r & \doteq & \sqrt{x_n^2 + y_n^2}, \\ f(r) & \doteq & 1 + k_c(1)r^2 + k_c(2)r^4 + k_c(5)r^6, \\ \mathbf{p}_d & = & \begin{bmatrix} f(r)x_n + 2k_c(3)x_ny_n + k_c(4)(r^2 + 2x_n^2) \\ f(r)x_n + 2k_c(4)x_ny_n + k_c(3)(r^2 + 2y_n^2) \end{bmatrix}. \end{array} \right. \quad (19)$$

Camera Calibration



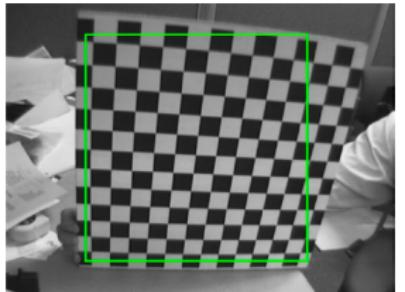
Multiple images \mathbf{I}_i of the calibration rig \mathbf{I}_o at different views τ_i are related by:

$$\mathbf{I}_i \circ (\tau_0 \circ \tau_i) = \mathbf{I}_o + \mathbf{E}_i, \quad i = 1, 2, \dots, N. \quad (20)$$

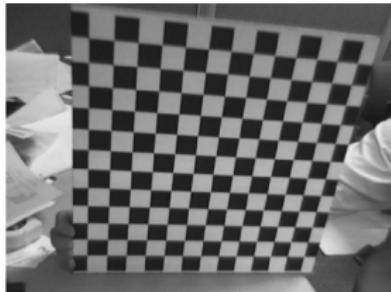
$$\min \sum_{i=1}^N \|\mathbf{L}_i\|_* + \|\mathbf{E}_i\|_1, \quad \text{s.t.} \quad \mathbf{I}_i \circ (\tau_0 \circ \tau_i) = \mathbf{L}_i + \mathbf{E}_i, \quad \mathbf{L}_i = \mathbf{L}_j. \quad (21)$$

$$\text{Or} \quad \min \|\mathbf{L}_c\|_* + \|\mathbf{L}_r\|_* + \lambda \|\mathbf{E}\|_1, \quad \text{s.t.} \quad \mathbf{I}_i \circ (\tau_0 \circ \tau_i) = \mathbf{L}_i + \mathbf{E}_i. \quad (22)$$

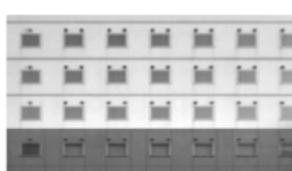
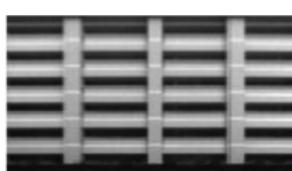
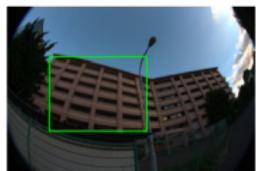
Calibration Results



(a) Input image with an initial window



(b) Lens distortion removed



Summary

A good idea deserves a good implementation.

The better the idea, the better should be the implementation.

Assignments

- Reading: Chapter 15.
- Programming Homework #4.